

Ralph Charles Evert IV

SENIOR AI LEADERSHIP

Enterprise AI Strategy · COEs · Op Models · Governance & Cost Control · Platform & Delivery

Miami Beach, FL · charlieevert@gmail.com · (215) 416-1737

[Portfolio](#)

[LinkedIn](#)

[GitHub](#)

[YouTube](#)

[Scholar](#)

[Books](#)



Applied AI leader who finds the problem worth solving, then solves it. Given a business area and a complaint rather than a specification, I scope the case, build the first version myself, then build and lead the forward-deployed team around it. Built the enterprise AI function from zero at three organizations — global CPG, federal, and financial services.

Consistently deliver against compressed timelines — a 12-month governance cycle closed in two, a three-month build shipped in three weeks. Currently manage a Director, Senior Managers, and Managers with full P&L, hiring, and margin ownership. US Army veteran. MBA and MS in Analytics.

How I build: evaluation before deployment, requirements before prompting, human review where confidence is low, measurable outcomes over demos.

EXPERIENCE

PwC

SEP 2025 — PRESENT · MIAMI, FL

Manager — Agentic AI Strategy & Implementation

Delivery director across two accounts: full ownership of the client relationship, the programs sold into them, and P&L. Lead onshore and offshore delivery centers mixing in-house and vendor talent — a Director in India, Senior Managers, and Managers reporting in on owned accounts.

ENTERPRISE STRATEGY & OPERATING MODEL

- Redesigned the digital and AI operating model for an F500 financial advisory leader stalled by pilot proliferation and disconnected governance — replacing it with domain-led delivery pods, a capability registry, and IT shifted toward platform maintenance.
- Advised the Chief Data & AI Officer and Chief Transformation Officer on budget allocation and value sequencing across the AI portfolio. Turned a manual operations team into a standing digital capability that tunes the system and prioritizes the gaps it surfaces.
- Run feature prioritization and change control independently across client roadmaps — pricing deviations as change orders so leaders see the tradeoff. One client's configuration request died once its cost was weighed against testing end users wanted.
- Built the value scorecard: workflow coverage, demand-driver mapping, and cost-per-contact targets, with satisfaction as a guardrail rather than a growth metric.

FUNCTION BUILDING & PEOPLE LEADERSHIP

- Founded and led the Agentic AI Customer Service Factory — converting reinvestment capacity into shippable products, workforce upskilling, and downstream revenue, built on a shared layer of reusable agents, MCP-exposed tools, and evaluation harnesses, so each new build started from assets rather than from scratch. Built and ran the matrixed team with a chief of staff; upskilled 45+ PMs and engineers. Own performance reviews for Managers and Senior Managers with full staffing, hiring, and margin authority.
- Authored delivery standards adopted across a 200+ practitioner practice, tying every build to reportable outcomes.

ORIGINATION & BUILD

- Uncovered a ~\$5M annual manual-labor pool inside a service center. Sized the case, won the funding, and built **SIGNAL** against it — 15-minute spike detection, semantic search and text-to-SQL over the transcript corpus, confidence-scored tagging with human review, and knowledge-gap routing. Sixteen weeks, five people; productized into a \$700K accelerator sold to a Fortune 500 bank.
- Built the AI security and evaluation regime for an enterprise voice and SMS platform — LLM-as-a-judge evals, simulated voice testing, and an adversarial telephony agent that attacks the system's own guardrails, surfacing prompt-injection paths before production.
- Pioneered experiential proposal responses — interactive digital experiences and a live AI agent evaluators could interrogate. Win rate moved from roughly 20% to 70% across ten proposals, now standard firmwide. Originated and delivered \$3M+ in AI sole-source engagements end to end.
- Initiated PwC's joint business agreement with xAI — sourced the relationship, built the value thesis, and ran it through to signature.

Deloitte

SEP 2023 – SEP 2025 · ARLINGTON, VA

Senior Consultant, AI & Engineering (promoted from Consultant)

Group product manager and technology lead across the federal health, defense, and HR technology portfolio.

Promoted within one year; selected for the firm's top performance tier.

- Ranked opportunities by value, momentum, and speed to impact across 20+ concurrent initiatives — recommending what to pursue, scale, or stop.
- Founded and led the firm's **Anthropic Alliance** from scratch — owning the roadmap, value thesis, and go-to-market. Grew to 250+ practitioners in six months, deploying 100 into product, advisory, and delivery. Deloitte's largest technology alliance of the year by funding. Originated the first joint delivery relationships with the FDA, AARP, the US Army, the US Navy, and the City of Miami.
- Group PM for the [Human Capital AI Suite](#) in a product operating model — roadmap, sequencing, versioning, and outcome ownership per deployment.
- Instrumented products for client-facing value reporting — volume processed, draft-to-publish cycle time, automated-versus-human handling. Products led generated \$1.35B+ in qualified pipeline and \$100M+ in realized P&L impact.
- Owned end-to-end delivery of [AARP](#)'s consumer AI platform serving 1M+ monthly users — the largest AWS Bedrock user base at initial deployment; sentiment 90%→30% negative, 10% renewal uplift. Built [Recruit360](#) for the US Army (12-week MVP), adapted to Coast Guard and Navy — 5x qualified leads, \$20M+ extension. Built [ReportGen](#) for FDA CDER One, cutting a two-week reporting cycle to three days.
- Led teams of up to 30, owning performance management end to end and advocating three engineers into promotion within a year each.

Johnson & Johnson (Kenvue)

SEP 2022 – SEP 2023 · SKILLMAN, NJ

Sr R&D Product Innovation Analyst (promoted from Data Engineer)

In-house digital transformation at a global CPG enterprise — decentralized, multi-market, multi-brand. Created the company's enterprise GenAI function. Promoted within six months.

- Ran the function as internal strategy consulting: scoped requirements with business units across brands and markets, built the opportunity portfolio, and sized **\$100M+ in automation opportunities** — implementing the highest-value ones personally.
- Automated the R&D and supply chain work behind the portfolio — raw-materials sourcing and traceability, formula cost calculation, supplier-compliance transcription, and packaging-compatibility checks — returning **10,000+ hours annually** and removing **20+ FTEs** of manual effort across procurement and product development. Built the R&D search platform that reverse-engineered competitor formulations from FDA and PubChem data.
- Built the enterprise GenAI function from scratch; the organization later hired a VP of Data Science & AI to lead it. Initiated vendor relationships with OpenAI and Pinecone directly, running comparative evaluation against Microsoft, AWS, and Google before the enterprise committed.
- Originated and built the organization's first enterprise GenAI platform, unprompted — productionized on Azure App Service and scaled to 22,000 users, delivered in partnership with Microsoft's Head of AI Product and OpenAI's Head of Sales. Delivered an **\$18M one-time cost avoidance** by replacing a planned commercial digitization program with a ~\$5K open-source OCR pipeline, then turned that corpus into the enterprise's first production RAG application.
- Led the GenAI squad building the **Enterprise Goal Setting Assistant**, serving all **22,000 employees** annually — directing UX/UI designers, software engineers, data engineers, and testers alongside stakeholders up to the **Chief People Officer and CIO**. The organization was missing its OKRs because employee goals did not tie back to them; the assistant closed that gap and cut the cost of the annual goal-setting cycle. The CIO named it a key accomplishment in [Authority Magazine](#).
- Co-authored enterprise digital governance policy with the CEO, CIO, SVP of Data, General Counsel, and CISO — preventing data leakage and liability exposure as employees adopted external AI tools — driving approval in two months against a typical 12-month cycle.

EARLIER ROLES

Graduate Research Assistant · Saint Joseph's University · Aug 2021 – Nov 2022 · Philadelphia, PA

Co-authored four peer-reviewed publications on NLP and econometric analysis. Built end-to-end Python pipelines with scikit-learn and NLTK — web scraping and data processing at scale, embeddings for semantic similarity, LDA topic modeling, sentiment analysis, clustering, and significance testing. Presented at conferences of 100+ attendees.

Data Analyst (Intern) · La Jolla International Gardens · Jun 2018 – Nov 2018 · San Diego, CA

Rent forecasting and revenue analysis for a 650+ unit residential community. Built Python pipelines for lead tracking.

Academic Programs Assistant · University of San Diego · Sep 2017 – May 2018 · San Diego, CA

Determined which students needed financial aid and built Python models to rank applicants on grades and other characteristics. Owned data entry, analysis, and financial reporting across the Salesforce applicant pipeline.

Airborne Infantryman (11B) · US Army · Jul 2015 – Nov 2016 · Fort Moore, GA

Operated under pressure with incomplete information. Honorable medical discharge for injury.

CAPABILITIES

STRATEGY

Business Case Development · Value Sizing · Portfolio Prioritization · Operating Model Design · Governance Forums · Scorecard Design · Executive Storytelling · P&L Ownership · Performance Management · Hiring · Vendor Management · C-Suite Advisory

DIGITAL & DATA

Enterprise AI Strategy · Agentic AI · RAG & Retrieval · LLM Evaluation · Product Management · Snowflake · Databricks · Palantir Foundry · Power BI

ENGINEERING

Python · FastAPI · React · Node.js · C#/.NET · SQL/NoSQL · Terraform · AWS Bedrock · Azure OpenAI · Claude Code · MCP servers & tool design · Reusable agent components · LangChain · CrewAI · Vector databases (Pinecone, Weaviate, OpenSearch) · CI/CD

EDUCATION

Saint Joseph's University · 2020–2022

MBA, Data Analytics (GPA 3.85)

MS, Business Intelligence + Analytics (GPA 3.95)

Temple University · 2017–2020

BBA, Accounting (GPA 3.4)

CERTIFICATIONS

AI & GENAI

Databricks Gen AI & LLM Industry Specializations — Healthcare & Life Sciences; Retail & Consumer Goods · 2025 · Salesforce Certified AI Associate · 2025 · DeepLearning.AI Multi AI Agent Systems with crewAI · 2024 · AWS Certified AI Practitioner · 2024 · Deloitte Advanced Prompt Engineer · 2024

CLOUD & DATA

AWS Certified Cloud Practitioner · 2024 · Microsoft Azure Fundamentals · 2024 · SAS Business Intelligence & Analytics · 2022 · Google Data Analytics Specialization · 2021 · Galvanize (Diligent) ACL Data Analytics Level 1 · 2020

ADDITIONAL

Google UX Design Specialization · 2021 · PA Real Estate License · 2020 · IRS Tax Preparation (VITA/TCE) · 2020 · PADI Open Water Diving

AWARDS & RECOGNITION

Amplify Impact Award — *PwC US Advisory* | **2026**

1 of 10 across a 36,000-person US Advisory practice, for AI-driven sales and delivery.

Outstanding Performance Award — *Deloitte* | **Dec 2024**

Led the 7-person team that took a consumer chatbot serving 1M+ users live.

STAR Award, Engineering — *Deloitte* | **Nov 2024**

Partner-selected, top 5 of 5,000+ practitioners.

Cavalry Award — *Deloitte* | **Oct 2024**

Quarterly top performer on Army Account

Applause Award, Federal Health — *Deloitte* | **Jul 2024**

Three agentic workflow demos presented to 50+ leaders.

Outstanding Performance Award — *Deloitte* | **Jul 2024**

Leadership behind the launch of the Anthropic Alliance.

Workday GenAI Ideathon Winner — *Deloitte* | **Apr 2024**

Led the winning team; productized as GoalAssist.

PUBLICATIONS, BOOKS & PATENT

Personal Profile Generator and Recommendation Engine — *US Patent 20250104393* | 2025

Fuel Costs and the EV Conversation: Public Sentiment Through Social Media — *Review of Business Information Systems* 26(1) | 2026

GenAI in Production: A Practical Guide to Evaluation and Optimization — *Book* | 2025

COVID-19 and Vaccine Dissemination on Airbnb's Popularity and Sentiment — *Int'l Journal of Web Based Communities* 21(1–2) | 2025

Comparative Study of LLM Evaluation Frameworks:

NLP vs. LLM-as-a-Judge — *IEEE SIEDS* | 2025

Whiskers and the Quantum Box: Quantum Computing for Ages 6–8 — *Book* | 2024

Ella and the Digital Adventure — *Book* | 2024

TEACHING & SERVICE

Deloitte x Anthropic Certification Program — co-created the skills and topic scope with Anthropic, who built the certification; basis for the public Claude Certified Architect credential.

PromptHub — video series on retrieval, evaluation, AI, and agent design.

Hire Heroes USA — Volunteer Advisory Board Member, 2023–Present. Mentored 50+ veterans on salary negotiation, upskilling, and interview prep.